

Repositories in Domain-Driven Design

Bridging Domain and Persistence in Go

Agenda

- 1 Why Repositories
- 2 Definition and Placement
- 3 Repository Styles
- 4 Aggregate Example
- 5 DAO vs Repository
- 6 Design Guidance
- 7 Testing Strategies
- 8 Key Takeaways

Why Repositories

Bridging Domain and Data

- Repositories mediate between domain objects and persistence, keeping the domain clean.
- They give the illusion of an in-memory collection while data lives in durable storage.
- Interfaces belong in the domain layer; infrastructure implements them behind the scenes.
- Promote expressive domain code such as `tenant.Activate()` without leaking database concerns.

Definition and Placement

Evans & Vernon on Repositories

- Evans: A Repository mediates between the domain and data-mapping layers, giving the illusion of an in-memory collection.
- Vernon reinforces the concept: expose repositories only for aggregate roots.
- Deliver domain objects in and out; hide records, SQL, and driver specifics.
- Domain layer relies on a stable contract; implementations can change via infrastructure adapters.

Repository Interface in Go

```
1 type TenantRepository interface {  
2     Add(ctx context.Context, tenant *Tenant) error  
3     ByID(ctx context.Context, id TenantID) (*Tenant, error)  
4     Remove(ctx context.Context, id TenantID) error  
5 }
```

- Methods reflect aggregate lifecycle operations rather than CRUD verbs alone.
- Keeps invariants within the aggregate while infrastructure handles serialization.

Repository Styles

Collection-Oriented Style

- Mimics an in-memory collection; Unit of Work tracks changes automatically.
- Common in ecosystems with rich ORMs (Java, .NET).
- Works best when the tooling offers change tracking and implicit transactions.
- Less common in Go, but useful to understand when collaborating across stacks.

Persistence-Oriented Style

```
1 type TenantRepository interface {  
2     Save(ctx context.Context, tenant *Tenant) error  
3     Delete(ctx context.Context, id TenantID) error  
4     ByID(ctx context.Context, id TenantID) (*Tenant, error)  
5 }
```

- Dominant approach in Go: explicit reads and writes per call.
- No hidden Unit of Work callers orchestrate state changes intentionally.
- Encourages clarity about I/O and transaction boundaries.

Aggregate Example

Tenant Aggregate Governing Users

```
1 type Tenant struct {  
2     ID      TenantID  
3     Name    string  
4     IsActive bool  
5     Users   []*User  
6 }  
7  
8 func (t *Tenant) Activate() { t.IsActive = true }  
9  
10 func (t *Tenant) RegisterUser(email string) *User {  
11     user := &User{Email: email, TenantID: t.ID}  
12     t.Users = append(t.Users, user)  
13     return user  
14 }
```

- Tenant acts as aggregate root, enforcing invariants for contained users.
- User data reaches persistence only through the tenant repository.

Application Service Coordinates Persistence

```
1 func (s *TenantService) ActivateTenant(ctx context.Context, id TenantID) error {  
2     tenant, err := s.repo.ByID(ctx, id)  
3     if err != nil {  
4         return err  
5     }  
6     tenant.Activate()  
7     return s.repo.Save(ctx, tenant)  
8 }
```

- Application layer owns transaction scope and orchestration.
- Repository remains focused on persisting aggregate state.

DAO vs Repository

Comparing DAO and Repository

Aspect	DAO	Repository
Focus	Tables / records	Aggregates / roots
Returns	DTOs or raw rows	Domain objects
Layer	Infrastructure	Domain interface, in-fra implementation
Concern	CRUD mechanics	Aggregate lifecycle & invariants

- Both can coexist: DAO encapsulates data access, repository composes domain-facing contract.
- Keeps the domain shielded from persistence technology churn.

Design Guidance

When to Use a Repository

- Provide one repository per aggregate root; avoid exposing internal entities.
- Shape methods around domain language *ActivateTenant*, *RegisterUser*, etc.
- Let repositories persist lifecycles; avoid mixing analytics or logging concerns.
- For cross-aggregate projections, create dedicated read models or query services.

Queries and Consistency

- Scope repository queries to identifiers or local keys within the aggregate boundary.
- Use query services/read models for reporting or multi-aggregate searches.
- Embrace CQRS-style separation when projections must remain eventually consistent.
- Keeps write-model repositories lean and focused on invariants.

Testing Strategies

Testing Repository Contracts

```
1 type InMemoryTenantRepo struct {  
2     store map[TenantID]*Tenant  
3 }  
4  
5 func (r *InMemoryTenantRepo) Save(ctx context.Context, t *Tenant) error {  
6     cp := *t  
7     r.store[t.ID] = &cp  
8     return nil  
9 }  
10  
11 func (r *InMemoryTenantRepo) ByID(ctx context.Context, id TenantID) (*Tenant,  
12     error) {  
13     return r.store[id], nil  
14 }
```

- In-memory fakes accelerate unit tests and validate domain logic.
- Integration tests with the real database prove mapping fidelity.
- Watch for aliasing or shared references in fakes to avoid optimistic tests.

Key Takeaways

- Repositories bridge domain behavior and persistence without leaking infrastructure.
- Prefer explicit, persistence-oriented methods in Go; know when collection-style applies.
- Keep repository interfaces small, aggregate-focused, and transaction-free.
- Separate read models for projections; test with both in-memory and database-backed implementations.
- A well-crafted repository sustains the domain between requests not merely a CRUD wrapper.